Text analysis

Jason G. Fleischer, Ph.D.

Asst. Teaching Professor Department of Cognitive Science, UC San Diego

<u>jfleischer@ucsd.edu</u>



@jasongfleischer

https://jgfleischer.com

Examples of questions that require text analysis

- 1. Did J.K. Rowling write <u>The Cuckoo's Calling under the pen name</u> Robert Galbraith?
- 2. What themes are common in 19th century literature?
- 3. Can we tell the difference between <u>tweets that come from Trump himself or a staffer</u>?
- 4. Is Hillary the most poisoned name in US History?
- 5. Can we visualize the <u>narrative structure of The Hobbit</u>?
- 6. Who has the biggest vocabulary in hip hop?
- 7. Is there a gender imbalance in who gets to say lines in the movies?

A whole bunch of text processing project ideas

- Spam detection (in text messages, on social media)
- Hate speech detection
- Predictive text (like gmail)
- Summarize articles (TL;DR)
- Summarize trends on social media... or in cutting edge NLP research?
- Plagiarism detection... in code projects?
- Document similarity... prevent the posting of duplicate questions on Campuswire?

Goal: Understand the basics of sentiment analysis

Today's example question: How has pop music

and TF-IDF

changed in the last five years?

What data would we need to answer this question?

How has pop music changed in the last five years?

Data: Lyrics to the most popular songs from each year

The data: Top songs from Feb music charts 2017-2021

2017: 152 songs

2018: 139 songs

2019: 127 songs

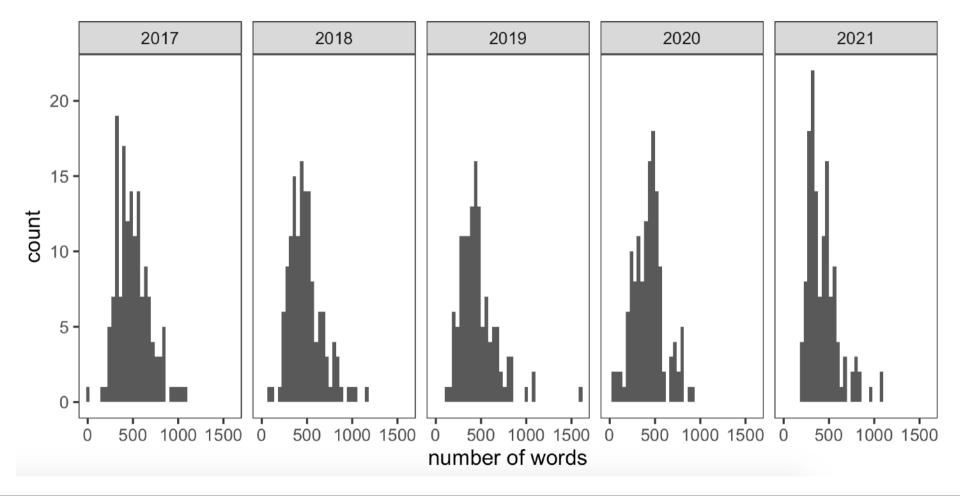
2020: 137 songs

2021: 134 songs

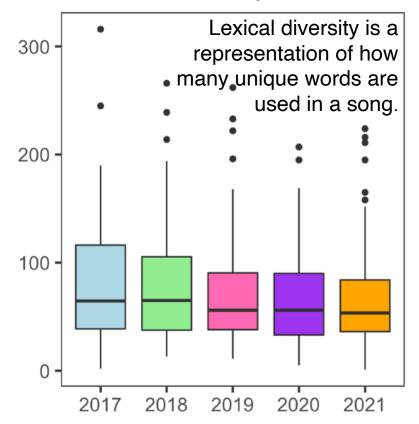
Song data from **Spotify**. Lyrics from **genius.com**

Questions we can ask...

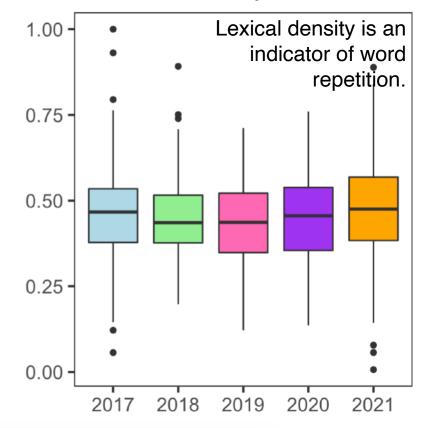
- Does the total number of words change over time?
- 2. Does uniqueness change over time?
- 3. Does the diversity or density change?
- 4. What words are most common?
- 5. What words are most unique to each year?
- 6. What sentiment do songs convey most frequently?
- 7. Has sentiment changed over time?
- 8. What are the sentiment of the #1 songs?
- 9. What words contribute to the sentiment of these #1 songs?
- 10....what about bigrams? N-grams?



Lexical Diversity



Lexical Density



Sentiment Analysis

Sentiment Analysis

Programmatically infer emotional content of text

text data text data

Break down
into a
individual or
combination of
words



compare to a sentiment lexicon: dataset containing words classified by their sentiment

Part of the "NRC" sentiment lexicon:

sentiment lexicon word <chr> <chr> <chr> abacus trust nrc abandon fear nrc abandon negative nrc abandon sadness nrc abandoned anger nrc abandoned fear nrc abandoned negative nrc abandoned sadness nrc abandonment anger nrc abandonment fear nrc ... with 27,304 more rows

When doing sentiment analysis...

token - a meaningful unit of text

- what you use for analysis
- tokenization takes corpus of text and splits it into tokens (words, bigrams, etc.)

Check out https://neptune.ai/

blog/tokenization-in-nlp

stop words - words not helpful for analysis

- extremely common words such as "a", "the", "of", "to"
- are typically removed from analysis

https://www.tidytextmining.com

When doing sentiment analysis...

stemming or lemmatization

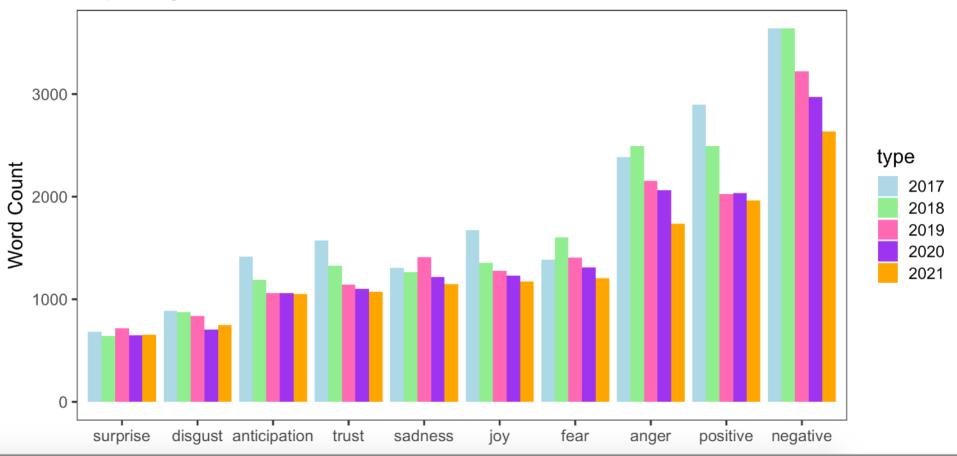
- Identifying the root for each token
- Jumping, jumped, jumps, jump all have the same root 'jump'
- Where things get tricky: jumper????
- Stemmers and lemmatizers do the same thing, stemmers are word focussed/cruder and lemmatizers try to take into account the context to correctly stem jumper differently depending on if its "I love my new jumper" (In the US jumper=a kind of dress, in the UK jumper=sweater) or its "Shane is a long jumper"

Check out https://www.datacamp.com/tutorial/stemming-lemmatization-python

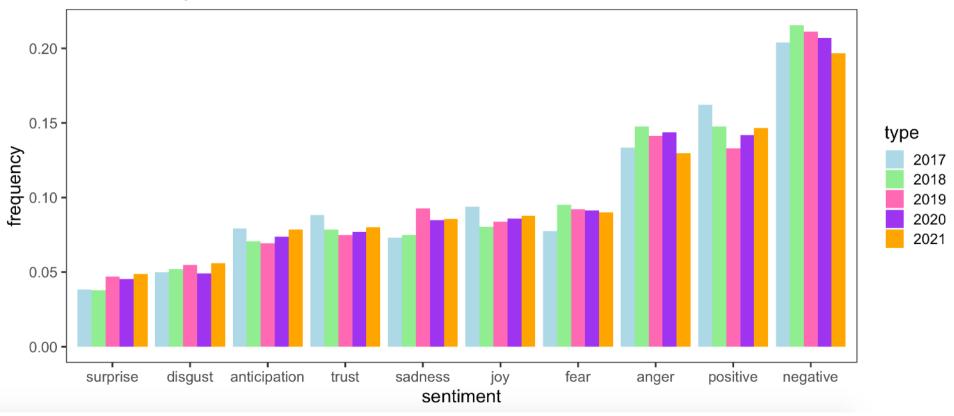
In text analysis, your choices matter:

- 1. How to tokenize?
- 2. Remove stop words? Remove common words?
- 3. Use stemming/lemmatization?
- 4. What sentiment lexicon to use?

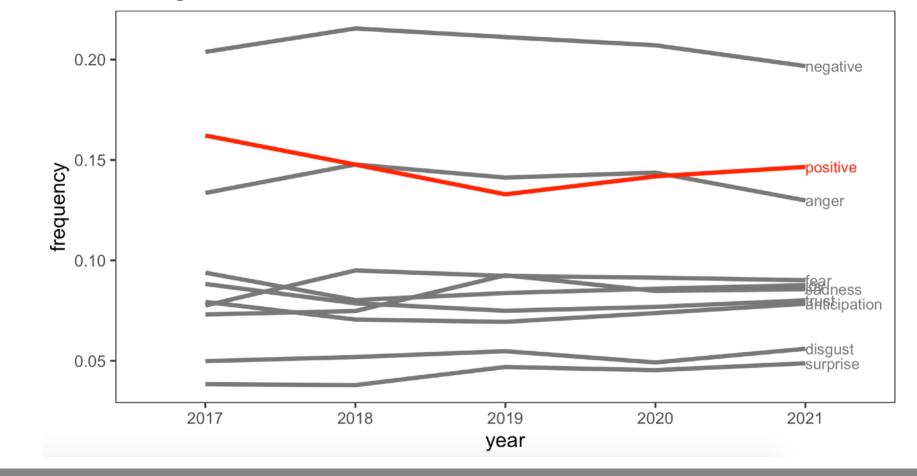




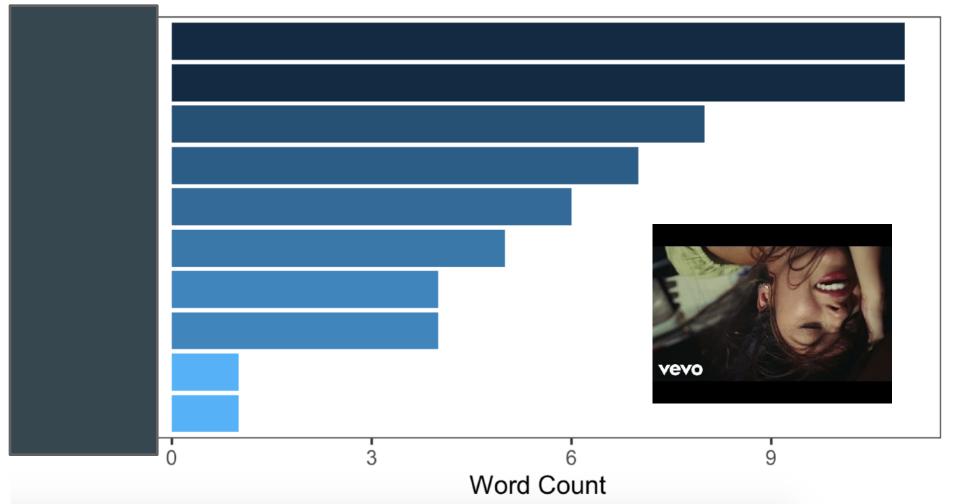
Sentiment by Year



Change in Sentiment over Time



Sentiment: Driver's License



I got my driver's license last week Just like we always talked about

'Cause you were so excited for me

To finally drive up to your house

But today I drove through the suburbs Crying 'cause you weren't around

And you're probably with that blonde girl Who always made me doubt

She's so much older than me

She's everything I'm insecure about

Yeah, today I drove through the suburbs

'Cause how could I ever love someone else? And I know we weren't perfect but I've never felt this way for no one

And I just can't imagine how you could be so okay now that I'm gone Guess you didn't mean what you wrote in that song about me 'Cause you said forever, now I drive alone past your street

And all my friends are tired

Of hearing how much I miss you, but I kinda feel sorry for them

'Cause they'll never know you the way that I do, yeah

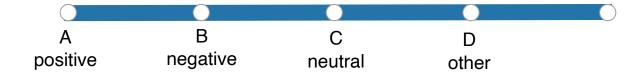
Today I drove through the suburbs And pictured I was driving home to you



Sentiment Limitations

How would you classify the sentiment of the following sentence?

"The idea behind the movie was great, but it could have been better"



"Dude!"

"So bad!"

The limitations of sentiment analysis: Context, homonyms, pragmatics, parts of speech

Mid, Shade

The limitations of sentiment analysis: Lexicon not up to the task, lexicon specific results

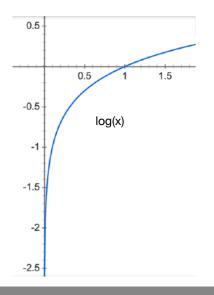
TF-IDF: Term frequency - Inverse document frequency

Given a corpus of documents, figure out which words are uniquely important to each document

TF-IDF:

Term Frequency - Inverse Document Frequency

Term Frequency (TF): how frequently a word occurs in a document Inverse document frequency (IDF): how important a word is to a document



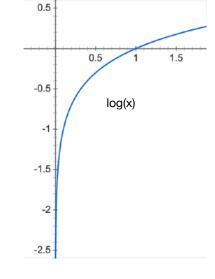
$$idf(ext{term}) = \ln \left(rac{n_{ ext{documents}}}{n_{ ext{documents containing term}}}
ight)$$

decreases the weight for commonly used words and increases the weight for words that are not used very much in a collection of documents

TF-IDF:

Term Frequency - Inverse Document Frequency the frequency of a term adjusted for how rarely it is used

$$w_{x,y} = tf_{x,y} \times log(\frac{n}{df_x})$$



Term x within document y

 $tf_{x,y}$ = frequency of x in y df_x = number of documents containing x

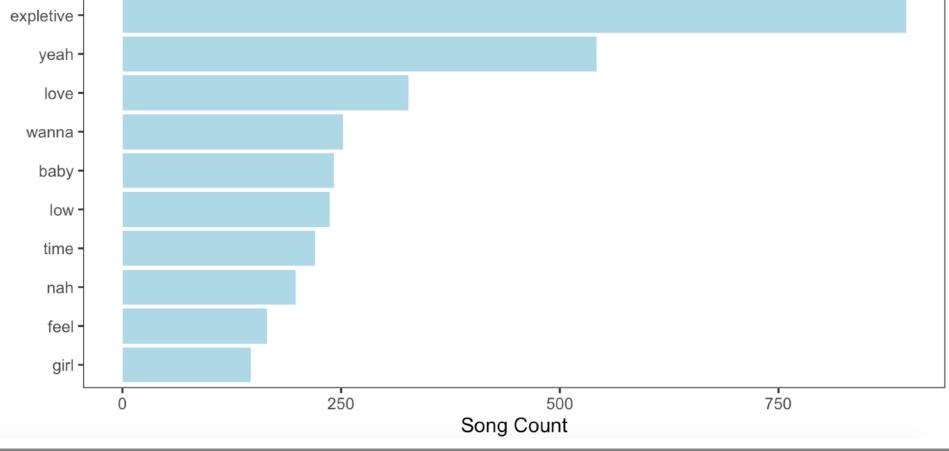
N = total number of documents

What words are the most unique to the lyrics of each year's top hits?

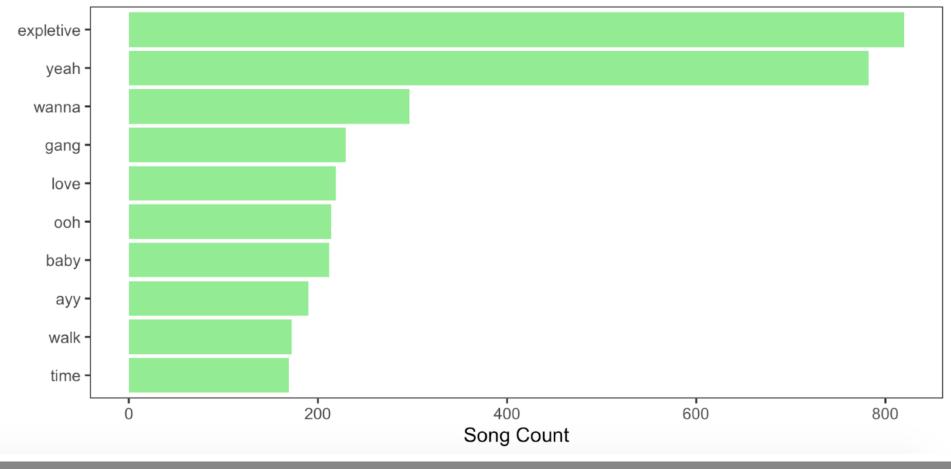
Goal: to use TF-IDF to find the important words for the content of each document by decreasing the weight for commonly used words and increasing the weight for words that are not used very much in a collection or corpus of documents

Calculating TF-IDF attempts to find the words that are important (i.e., common) in a text, but not *too* common

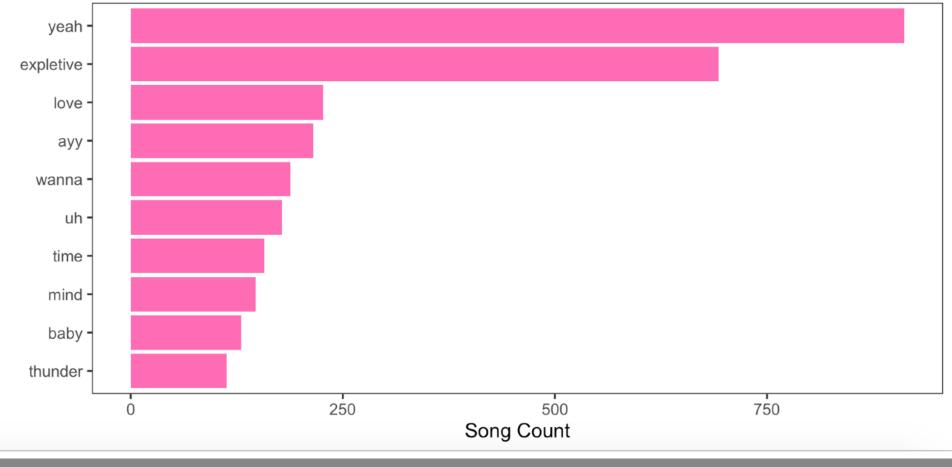
Most Frequently Used Words in top 200 songs (2017)



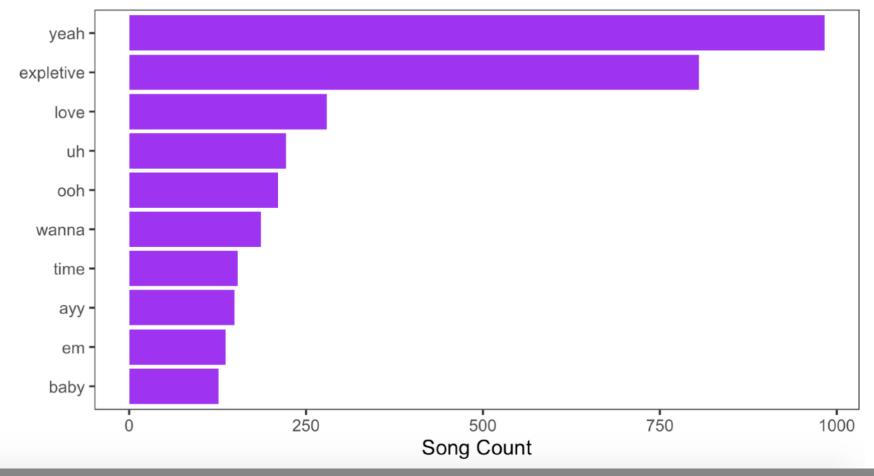
Most Frequently Used Words in top 200 songs (2018)



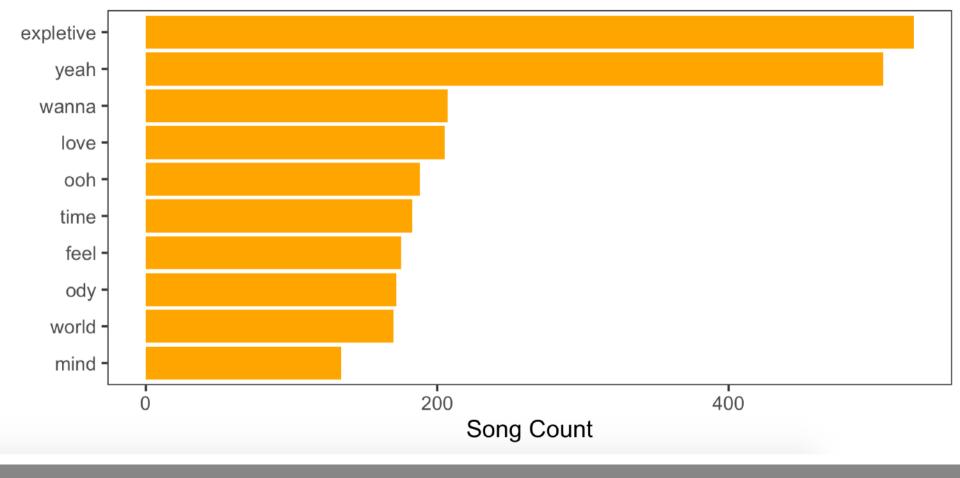




Most Frequently Used Words in top 200 songs (2020)



Most Frequently Used Words in top 200 songs (2021)







Term
Frequency
can only tell

us so much.... 2017



2019

2018



2020



TF-IDF: Term Frequency - Inverse Document Frequency

Term Frequency (TF): how frequently a word occurs in a document

Inverse document frequency (IDF): intended to measure how important a word is to a document

decreases the weight for commonly used words and increases the weight for words that are not used very much in a collection of documents

$$idf(ext{term}) = \ln \left(rac{n_{ ext{documents}}}{n_{ ext{documents containing term}}}
ight)$$

TF-IDF:

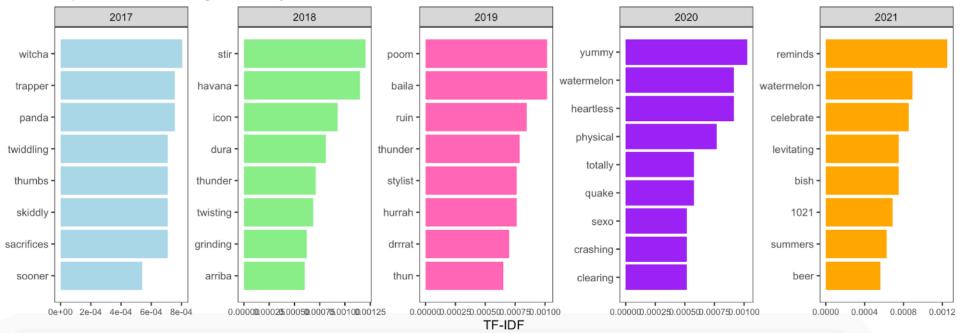
Term Frequency - Inverse Document Frequency the frequency of a term adjusted for how rarely it is used

$$w_{x,y} = tf_{x,y} \times log(\frac{iv}{df_x})$$



 $tf_{x,y}$ = frequency of x in y df_x = number of documents containing x N = total number of documents

Important Words using TF-IDF by Year



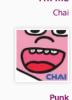
Flight of the Conchords

Father and Son

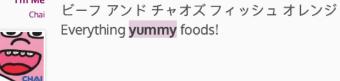
Driving round in the car I even let you drive Eating dinner from a can dad's **yummy** can delight 2019

Just you and I





2019

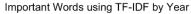


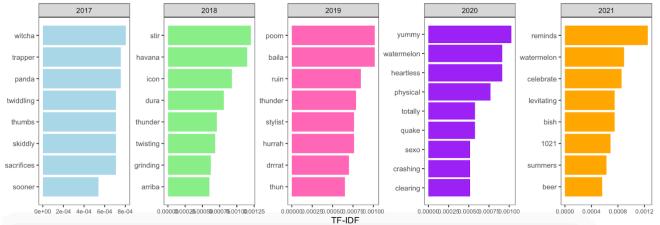


Sex Talk 2019 He eat from the back to the front Bitch I got green like I'm Buttercup (ah) He say I got that lil **yummy** yum (ah ah)



What can you conclude from this TF-IDF plot?





A No words overlap across the years in these data



B 'reminds' and 'watermelon' are the most unique words to the 2021 data

C 'watermelon' is the most common word in this dataset

D A-C (all of the above)

E None of the above

https://forms.gle/uD8cVTUUzAsxcVvKA



Questions we can ask...

- Does the total number of words change over time?
- 2. Does uniqueness change over time?
- 3. Does the diversity or density change?
- 4. What words are most common?
- 5. What words are most unique to each year?
- 6. What sentiment do songs convey most frequently?
- 7. Has sentiment changed over time?
- 8. What are the sentiment of the #1 songs?
- 9. What words contribute to the sentiment of these #1 songs?
- 10....what about bigrams? N-grams?

EDA

TF-IDF

Sentiment Analysis

Text analysis

The big picture

- Identify the problem you are trying to solve... what you want to answer suggests the approaches you will use below
- Clean and preprocess the data
- Linguistic transformations
 - Tokenization, lemmatization, stemming, remove stop words
 - [optional] part of speech tagging, named entity recognition
- EDA

- [optional] Calculate measurements and statistics
 - e.g., diversity, density, sentiment, TF-IDF
- [optional] Transform representation
 - Bag of words, vector embeddings
- [optional] Train a ML system or write an algorithm to accomplish a task given the representation
 - e.g., detect spam